

Introduction to Big Data

Concepts, Solutions, and Tools

Mahmoud Parsian, Ph.D. in Computer Science

Agenda

1. What is Big Data? Why now?
2. The 5 V's, file formats, and types of analytics
3. Foundational concepts: scaling, CAP, ACID vs. BASE, partitioning, modeling, governance & security
4. The Big Data life cycle — batch, real-time, and orchestration
5. Solutions: Hadoop & MapReduce, Apache Spark, Amazon Athena, and the cloud platform landscape
6. Glossary, decision checklist, case study

What Is Big Data?

Big Data is an umbrella term for **non-traditional strategies and technologies** needed to:

- **Gather** data
- **Organize** data
- **Process** and **Analyze** data
- **Extract insight** from datasets too large, fast-moving, or varied for a single machine or traditional software tools to handle



What Is Big Data? (cont'd)

- Large, complex, and constantly changing datasets
- **Cannot** be stored, managed, or processed adequately by traditional software and tools (e.g., a single RDBMS on one server)
- Datasets of billions of elements — measured in **terabytes** and **petabytes**

Honest answer: there is no single, universally agreed definition — projects, vendors, and practitioners each use the term somewhat differently. What follows is the working definition this course uses.

Why Big Data, Why Now?

Explosive data growth driven by:

- **Social networking** — Twitter/X, Facebook, ...
- **Media** — Netflix, streaming, cable TV, ...
- **Commerce** — Amazon, eBay, ...
- **Health & genomics** — data growing exponentially
- Every year the world creates more **zettabytes** of data


Managing this data well is key to business expansion — in retail, finance, media, and publishing.

Examples of Big Data

- Credit card transactions
- Billions of documents indexed for a search engine
- Logs generated by web servers
- DNA & genomics data
- Twitter/X feeds, Facebook messages
- Medical records, COVID data

Big Data's Key Challenges

- **Storage** — petabytes of data
- **Processing** — must be fast and effective
- **Diversity of data** — structured, semi-structured, multi-structured, unstructured
- **Cost** — increased cost to manage and process large datasets



The 5 V's of Big Data

V	Meaning
Volume	Size of the data (GB → TB → PB)
Velocity	Speed at which data is generated / must be processed
Variety	Structured, semi-structured, unstructured sources
Veracity	Quality / trustworthiness of the data
Value	Whether the data is actually worth the cost to process

Some frameworks also call out **Complexity** as a 6th dimension (interrelated, multi-source data that's hard to link and reconcile).

Volume

- Datasets now frequently exceed:
 - **1 TB = 1,024 GB**
 - **1 PB = 1,024 TB**
- Too large to process with a regular laptop/desktop CPU
→ requires **cluster computing** + a **distributed file system**

Example — credit cards: 369 billion purchase transactions worldwide in 2018.

Example — DNA samples: 1,000,000 samples × 5,000,000 records each
= 5 trillion data points.

Velocity

- The speed at which data is generated and must be ingested/processed
- High-velocity data needs **distributed / streaming** processing techniques, not batch-only tools

Examples: text messages, Twitter/X messages, Facebook posts, sensor telemetry, credit-card swipes.

Variety

Big Data comes from many kinds of sources, generally one of:

- **Structured** — rows/named columns (CSV, Parquet, RDBMS tables)
- **Semi-structured** — JSON, XML
- **Unstructured** — text files, log files, audio, video, images

Some also call out **multi-structured** data — a mix of the above in one source (e.g., a web log with structured fields plus a free-text message). High-variety data like this often needs specialized processing and algorithms.

Big Data File Formats: Row vs. Columnar

- **Row-oriented** (CSV, JSON, Avro) — stores each record contiguously; fast to write, fast to read a whole row (OLTP-style access)
- **Column-oriented** (Parquet, ORC) — stores each column contiguously
 - Fast for analytical scans that touch few columns over many rows
 - Compresses better — similar values sit next to each other
 - What Spark, Athena, and most Big Data query engines prefer

Rule of thumb: row formats for streaming/transactional writes, columnar formats for analytical reads.

Veracity

- Refers to the **quality/trustworthiness** of the data
- Low veracity → high percentage of **noise** → wrong conclusions
- High veracity data (e.g., a controlled medical trial) is reliable enough to drive real decisions

Garbage in, garbage out applies at Big Data scale too — bad data quality doesn't average out just because you have more of it.

Value

- Volume, Velocity, Variety, and Veracity describe the *data itself*; **Value** asks whether processing it is actually **worth it**
- Not all data is worth the cost of collecting, storing, and processing — weigh cost against benefit before building a pipeline
- Value is usually unlocked only *after* the full life cycle: raw data → cleaned/aggregated → analyzed → **decision or insight**

Data is only "Big Data" in a useful sense if someone can turn it into value.

Types of Big Data Analytics

Four questions analytics can answer, in increasing sophistication:

Type	Question	Example
Descriptive	What happened?	Monthly sales dashboard
Diagnostic	Why did it happen?	Root-cause of a sales dip
Predictive	What will happen?	Churn / demand forecasting
Prescriptive	What should we do?	Recommend the next best action

Big Data enables all four at scale — more data gives more signal for the harder predictive/prescriptive questions.

Roles in the Big Data World

- **Data Engineer** — builds and operates the pipelines (ingest, storage, compute) covered in this deck: Spark, Kafka, Airflow, ...
- **Data Analyst / BI Analyst** — descriptive & diagnostic analytics: dashboards, reports
- **Data Scientist** — predictive & prescriptive analytics: statistical models, ML
- **ML Engineer** — productionizes and scales the models data scientists build

These roles overlap heavily in practice — Big Data infrastructure is what makes all of them possible at scale.

Foundational Concept: Vertical vs. Horizontal Scaling

- **Vertical scaling (scale up)** — add more CPU/RAM/disk to *one* machine. Simple, but has a hard ceiling and a single point of failure.
- **Horizontal scaling (scale out)** — add *more machines* to a cluster and spread the work. This is the Big Data approach:
 - 10s, 100s, or 1,000s of commodity servers
 - **Shared-nothing architecture** — each node owns its own CPU/memory/disk; nodes coordinate over the network instead of sharing physical resources
 - Fault tolerance via **replication**, not exotic hardware

Foundational Concept: CAP Theorem

A distributed data store can only guarantee **two of three** properties at the same time:

- **C**onsistency — every read sees the latest write
- **A**vailability — every request gets a (non-error) response
- **P**artition tolerance — the system keeps working despite network partitions between nodes

Because network partitions *will* happen at scale, real systems pick **CP** (e.g., HBase — consistent, may reject requests during a partition) or **AP** (e.g., Cassandra, DynamoDB — always answers, may briefly serve stale data).

Foundational Concept: ACID vs. BASE

	ACID (traditional RDBMS)	BASE (many NoSQL / Big Data stores)
Stands for	Atomicity, Consistency, Isolation, Durability	Basically Available, Soft state, Eventual consistency
Consistency	Strong, immediate	Eventual — replicas converge over time
Optimized for	Correctness on a single/small cluster	Availability & scale across many nodes

Big Data systems often trade ACID for BASE's availability and horizontal scale — a direct consequence of the CAP theorem.

Foundational Concept: Partitioning & Sharding

Splitting a dataset so no single node/file holds it all.

Partitioning splits data within a system (e.g., a table into files); **sharding** spreads it across separate machines/instances.

- **Hash** — `hash(key) % N` picks the partition; even load, but range queries must hit every partition
- **Range** — contiguous key ranges (e.g., by date); great for range queries, risks hot spots
- **List / value** — one partition per discrete value (e.g., by `continent`) — what Athena/S3 partitioning uses (covered later)

Choosing a partition key is a **modeling decision**, driven by how the data will be queried.

Foundational Concept: Data Lake vs. Warehouse vs. Lakehouse

- **Data lake** — large repository of raw, often unstructured or semi-structured data, stored as-is (schema-on-read; e.g., raw files in S3/HDFS)
- **Data warehouse** — cleaned, integrated, well-ordered data, optimized for reporting/analysis (schema-on-write)
- **Data lakehouse** (modern evolution) — adds warehouse-style structure, ACID transactions, and schema enforcement *on top of* cheap lake storage (e.g., Delta Lake, Apache Iceberg, Apache Hudi) — you get one copy of data serving both raw and curated needs

Foundational Concept: Modeling for Big Data

- **Relational modeling** — normalize to reduce redundancy; joins computed at query time; optimized for consistent, small updates
- **Big Data modeling** often **denormalizes** on purpose:
 - Wide, nested tables (structs/arrays) avoid expensive distributed joins
 - **Star / snowflake schemas** — a fact table plus dimension tables, still common for warehouse-style analytics
 - Trades storage (cheap) for avoiding compute-heavy joins (expensive at cluster scale)

Model for how the data will be read, not just how it's produced.

Foundational Concept: Governance, Security & Privacy

- **Access control** — who can read/write which data (e.g., Apache Ranger, AWS IAM policies)
- **Encryption** — at rest (disk/S3) and in transit (TLS)
- **Data lineage & cataloging** — where data came from, what transformed it (e.g., AWS Glue Data Catalog, Hive Metastore)
- **Compliance & ethics** — regulatory requirements on sensitive data (GDPR, HIPAA for health/genomics data, CCPA), and awareness that models trained on biased data perpetuate that bias at scale
- **Auditing** — who accessed what, and when

Big Data's scale makes governance *harder* and *more necessary* — more data, more copies, more places for it to leak.

Where Do We Store Large Datasets?


- **Amazon S3** (object storage — the modern default)
- **Distributed file systems:**
 - Hadoop Distributed File System (**HDFS**)
 - Google File System (GFS)
 - NFS (Network File System)
- **NoSQL storage alternatives:**
 - Document stores — Apache CouchDB, MongoDB
 - Graph stores — Neo4j
 - Key-value stores — Apache Cassandra, Riak
 - Tabular/column stores — Apache HBase

Why Distributed Storage? An NFS Example

100 file servers, 200 TB each → **20,000 TB**
total raw space.

Replication factor	Usable space
1	20,000 TB
2	10,000 TB
4	5,000 TB

Replication costs money — it trades raw capacity for fault tolerance and read parallelism.



The Big Data Life Cycle

1. **Create** data from many sources
2. **Ingest** data into the system
3. **Persist** the data in storage
4. **Compute & Analyze** the data
5. **Visualize & Present** the results

Each stage typically uses different, purpose-built tools.

1–2. Creating & Ingesting Data

Data ingestion = taking raw data and adding it to the system.

Complexity depends on source format/quality and how far it is from the desired end state. Dedicated ingestion tools:

- **Apache Kafka** — high-throughput streaming ingestion; today's default choice
- **Apache Sqoop** — bulk-load from relational DBs (*retired to the Apache Attic; largely replaced by managed ELT tools*)
- **Apache Flume, Apache Chukwa** — older log/event collectors, still seen in legacy Hadoop deployments
- **Gobblin**

3. Persisting Data in Storage

Ingestion hands data to a storage layer that must handle volume, availability, and distributed access reliably:

- Amazon S3
- Hadoop's HDFS
- Google File System
- NFS

4. Computing & Analyzing Data — Batch


- Split work into pieces → schedule on machines → shuffle intermediate results → reduce → assemble final result
- Steps: *split, map, shuffle, reduce, assemble* — a distributed MapReduce algorithm
- Best for very large datasets, less time-sensitive

4. Computing & Analyzing Data — Real-Time

- Processes individual items continuously as they arrive
- Pairs with **in-memory computing** to avoid disk round-trips, and scales **elastically** as the stream's rate changes
- Apache Storm, Apache Flink, Apache Spark (Structured Streaming)

Batch vs. Real-Time, Compared

	Batch	Real-Time / Streaming
Trigger	Periodic or one-time	Continuous, 24/7
Latency	Minutes to hours	Milliseconds to seconds
Fits	Large, less urgent data	High-velocity, time-sensitive data



Lambda vs. Kappa Architecture

Two standard patterns for combining batch + real-time processing:

- **Lambda architecture** — run *both* a batch layer (accurate, slow) and a speed layer (approximate, fast) in parallel; merge results at query time
 -  correctness + low latency  two code paths to maintain
- **Kappa architecture** — treat *everything* as a stream; a "batch" is just replaying the stream from the beginning
 -  one code path  needs a stream engine that can also do full reprocessing (e.g., Kafka + Spark Structured Streaming)

5. Visualizing the Results

Final results are often saved to:

- ElasticSearch
- Relational databases
- Snowflake

Visualization / exploration tools:

- Jupyter Notebook
- Apache Zeppelin
- BI tools connected via JDBC/ODBC (to Athena, Snowflake, ...)

Orchestrating the Pipeline

The life-cycle stages (ingest → persist → compute → visualize) don't run themselves — a **workflow orchestrator** schedules and chains them as a **DAG** (directed acyclic graph) of steps, handling retries, dependencies, and failures:

- **Apache Airflow** — Python-defined DAGs, today's de facto standard
- **Apache Oozie** — Hadoop-native workflow scheduler
- Cloud-managed equivalents: AWS Step Functions, Google Cloud Composer, Azure Data Factory

Big Data Solutions

MapReduce → Hadoop → Spark → Amazon Athena

The Landscape of Solutions

Technologies that store, process, and analyze large quantities of data:

- **MapReduce** — the model/paradigm
- **Apache Hadoop** — the first popular open-source implementation
- **Apache Spark** — the modern, in-memory superset of MapReduce
- **Amazon Athena, Google BigQuery, Snowflake** — serverless SQL-on-object-storage
- **Apache Tez** — another MapReduce-style execution engine

Cloud Big Data Platforms

Capability	AWS	Azure	Google Cloud
Object storage	S3	Blob Storage	Cloud Storage
Managed Hadoop/Spark	EMR	HDInsight	Dataproc
Serverless SQL	Athena	Synapse Serverless	BigQuery
Managed Spark SaaS	Databricks	Databricks	Databricks

This course focuses on **AWS** (S3, EMR, Athena) — the concepts transfer directly to the other clouds' equivalents.

The Paper That Started It All (2004)

"MapReduce: Simplified Data Processing on Large Clusters" —

Jeffrey Dean & Sanjay Ghemawat, Google, OSDI'04.

- Google built MapReduce to index billions of web documents
- Google did **not** release the code publicly
- **Apache Hadoop** was built as an open-source implementation of the paper's ideas, becoming an Apache project in 2006

MAPREDUCE: SIMPLIFIED DATA PROCESSING ON LARGE CLUSTERS

by Jeffrey Dean and Sanjay Ghemawat

Abstract

MapReduce is a programming model and an associated implementation for processing and generating large datasets that is amenable to a broad variety of real-world tasks. Users specify the computation in terms of a *map* and a *reduce* function, and the underlying runtime system automatically parallelizes the computation across large-scale clusters of machines, handles machine failures, and schedules inter-machine communication to make efficient use of the network and disks. Programmers find the system easy to use: more than ten thousand distinct MapReduce programs have been implemented internally at Google over the past four years, and an average of one hundred thousand MapReduce jobs are executed on Google's clusters every day, processing a total of more than twenty petabytes of data per day.

MapReduce: The Model

MapReduce is a model/paradigm, not a product — a way to:

1. **Partition** data into smaller chunks
2. **Parallelize** transformations on the chunked data

```
map(input)      -> {(key1, value1), (key2, value2), ...}  
reduce(key, [v1, v2, ...]) -> {(key, final_value), ...}
```

- **map** splits the problem and computes over the pieces in parallel
- **shuffle** aligns intermediate results by key
- **reduce** combines the values for each key into one output

Full derivation and worked examples: see

slides/mapreduce/mapreduce_by_Mahmoud_Parsian/ .

Apache Hadoop

- Framework to efficiently **store** (HDFS) and **process** (MapReduce) datasets from gigabytes to petabytes
- **HDFS** — distributed file system for storage
- **YARN** (Yet Another Resource Negotiator) — cluster resource management/scheduling
- **MapReduce** — the batch computation engine

Hadoop tools and services work together as the **Hadoop ecosystem**.

Deep dive on HDFS: see

[06_introduction_to_hdfs.md](#) .

The Hadoop Ecosystem

Component	Role
HDFS	Distributed file system for storage
MapReduce	Parallel processing framework
ZooKeeper	Coordination framework
Hive	SQL-like interface for querying
Pig	High-level scripting for batch/ETL
HBase	Tabular/column, non-relational, distributed DB
Flume	Distributed service for collecting/moving data

Problems With Hadoop

- **Complex** — lots of low-level code, low-level API
- Uses **disk I/O** extensively; doesn't use RAM
(RAM is ~100× faster than disk I/O)
- Historically **text-file-centric**
- **No multi-language support** (Java-only, originally)

Hadoop's problems → motivated Apache Spark.

Apache Spark

Spark is a **unified engine for large-scale data analytics**:


- Multi-language engine — Python, Java, Scala, SQL, R
- Runs on a single node or a cluster
- Uses **memory/RAM first, then disk**
- Parallel processing — a **superset of MapReduce**
- **Data partitioning** for parallelism, built-in scalability
- Data abstractions: **RDDs** and **DataFrames**
- Enables **SQL access** to Big Data (Spark SQL)

Spark's Unified Stack

Spark Core, with higher-level libraries built on top:

- **Spark SQL** — interactive queries over structured data
- **Structured Streaming** — real-time analytics
- **MLlib** — machine learning
- **GraphX** — graph processing


All accessible from **Python, Scala, Java, or R**.



Spark Cluster Architecture

- **Driver program** — runs your `SparkContext`, plans the job
- **Cluster manager** — allocates resources across the cluster
- **Worker nodes** — each runs an **Executor** that runs **Tasks** in parallel

The driver splits work into tasks and ships them to executors, which process data partitions in parallel and report results back.



PySpark Example: Word Count

```
# word_frequencies : {(word1, freq1), (word2, freq2), ...}
input_path = "s3://my_bucket/project23/"
# spark: an instance of SparkSession
word_frequencies = (
    spark.sparkContext.textFile(input_path)
        .flatMap(lambda x: x.split(' '))
        .map(lambda x: (x, 1))
        .reduceByKey(lambda a, b: a + b)
)
```

Same MapReduce *model* as before — expressed in ~4 lines instead of a full Java MapReduce job.

More worked examples: [slides/mapreduce/.../03_word_count_in_python.md](#)

Hadoop vs. Spark

	Hadoop	Spark
Programs	Long, complex, low-level API	Short, simple, high-level API
Languages	Java	Python, Java, Scala, SQL
Speed	Slow — disk I/O bound	Fast — uses RAM as much as possible
Processing	Parallel (MapReduce)	Parallel (superset of MapReduce)
Maturity	Established, older	De facto standard for new work

Spark is generally the better default choice today for new Big Data processing pipelines.

SQL Access to Big Data

Multiple technologies now offer SQL-like access on top of Big Data storage:

- **Amazon Athena (JDBC), Snowflake (JDBC), Impala, Hive, Trino/Presto**
- **Spark SQL** — SQL access to Spark DataFrames

Used for: data-science sandboxes, ad-hoc analytics, search indexes over multi-structured data, data visualization pipelines.

Amazon Athena

Serverless SQL on data already sitting in S3

What Is Amazon Athena?

- **Interactive, serverless SQL query service** on AWS
- Analyzes data **directly in Amazon S3** — no data movement required
- **Serverless:**
 - No cluster to set up or manage
 - **Pay-per-query** — you pay only for data scanned (~\$5 per TB scanned)
- Built on **Trino/Presto** (originated as Facebook's PrestoDB; Athena's newer engine versions run on Trino)
- Supports CSV, JSON, ORC, Avro, **Parquet**, and more
- Client access via a web console or **JDBC/ODBC**

Athena Concepts

- **Database** — a logical grouping (catalog) of tables
- **Table** — metadata describing your data; the actual data lives in S3. Deleting a table does **not** delete the underlying S3 data
- **SerDe** (Serializer/Deserializer) — tells Athena/Hive how to parse a given data format
- Optimized for **columnar formats** (Parquet, ORC); data can also be compressed (e.g., GZIP) to cut scan cost

Amazon Athena: Use Cases

- Ad-hoc querying of logs, genomics/DNA data, or social data in S3 when latency isn't critical
- Going **fully serverless**: API Gateway → Lambda → Athena → S3, with no persistent compute to manage
- Mixing AWS services: run an on-demand **EMR** job to produce results in S3, then query them with Athena
- **Note**: no traditional multi-statement transactions (Athena does support single-table ACID operations on Apache Iceberg tables)

Full tutorial with partitioning walkthrough:

[slides/amazon_athena/](#).

Considering a Big Data Solution?

Ask before reaching for Hadoop/Spark/Athena — one question per V:


- Does the problem require handling high **volume**?
- Does it require handling high **velocity**?
- Does it require handling **variety** of data types?
- Can you trust the data's **veracity** enough to act on it?
- Will the **value** extracted justify the cost of the solution?
- Can the solution handle the **complexity**, and can it be **optimized** for cost and speed?

If the answer to all of these is "no," a traditional RDBMS may still be the right tool.

Where Do NoSQL Stores Fit In?

NoSQL databases trade strict relational guarantees for flexibility and distributed-first architecture — a natural fit for Big Data:

- **Key-Value** — Cassandra, Riak
- **Column-Family** — HBase
- **Document** — MongoDB, CouchDB
- **Graph** — Neo4j
- **SQL-on-Big-Data** — Spark SQL, Amazon Athena, BigQuery, Snowflake



Big Data Glossary

- **Batch processing** — process data in large, non-time-sensitive sets; start the job, get results later
- **Cluster computing** — pooling multiple machines' resources under a management layer that coordinates work and communication
- **Data lake** — large repository of raw, relatively unstructured collected data
- **Data warehouse** — cleaned, integrated, well-ordered data for analysis/reporting
- **ETL** — Extract, Transform, Load: preparing raw data for use

Big Data Glossary (cont'd)

- **Data mining** — finding patterns in large datasets
- **In-memory computing** — keeping working data and intermediate results in cluster memory instead of writing to disk between steps (Spark's core speed advantage over Hadoop MapReduce)
- **Machine learning** — systems that learn/improve from data via predictive and statistical algorithms
- **NoSQL** — databases designed outside the relational model, often distributed-first
- **Stream processing** — computing over individual items continuously as they arrive

Case Study: Social Media Analytics

Using people's online history — purchases, searches — to build a picture of attitude toward a product or brand.

These outputs feed downstream analyses such as:

- Customer satisfaction
- Churn prediction
- Financial performance
- Stock performance

Summary

- **Big Data** = data too large/fast/varied for traditional tools — characterized by the **5 V's** (Volume, Velocity, Variety, Veracity, Value)
- Handling it requires **horizontal scaling**: cluster computing + distributed storage, with trade-offs described by **CAP** and **ACID vs. BASE**
- **MapReduce** is the foundational model; **Hadoop** was the first popular implementation; **Spark** is today's in-memory, multi-language default

Summary (cont'd)

- **Amazon Athena** (and peers like BigQuery/Snowflake) offer **serverless SQL** directly on data in object storage — one of several equivalent options across AWS, Azure, and Google Cloud
- Modeling, partitioning, governance, and orchestration turn raw storage + compute into a **usable, safe, maintainable** pipeline
- Choice of tool follows from the life-cycle stage (ingest, persist, compute, visualize) and the workload's V's

Questions?

Mahmoud Parsian

Ph.D. in Computer Science