

Introduction to HDFS

Hadoop Distributed File System

Mahmoud Parsian

Ph.D. in Computer Science

What's Hadoop?

An open-source framework that:


- Efficiently **stores** large datasets, from gigabytes to petabytes.
- Efficiently **processes** those datasets, by implementing the MapReduce paradigm (`map()` , `reduce()` , optional `combine()`).

HDFS — Hadoop Distributed File System — is the piece that handles storage: a distributed file system for large datasets running on commodity hardware.

What's HDFS?

- **Fault tolerant, scalable**, and easy to expand.
- The **primary distributed storage** for Hadoop applications.
- Provides interfaces that let applications move their computation *closer to the data* — the same "close to the data" scheduling principle from `04_word_count_in_mapreduce.md`.

Where HDFS Fits in the Cluster



Where HDFS Fits in the Cluster (Continued)


Applications submit work to a Cluster Manager, which dispatches it to executors on worker nodes. HDFS is the storage layer those workers read from and write to.

Sizing Example

A 15-node cluster: 1 Master (stores metadata only, no actual data) + 14 worker nodes, each able to store 100TB:

```
total usable storage = 14 x 100TB = 1400TB
```

HDFS Architecture



Two Kinds of Nodes

- **NameNode** (master) — the heart of the filesystem. Maintains and manages metadata only: which blocks make up a file, and which DataNodes those blocks live on. Stores no actual data.
- **DataNode** (worker) — stores the actual data, and runs mappers and reducers.

A **Secondary NameNode** backs up the namespace (metadata), and DataNodes send heartbeats so the NameNode can direct replication and load balancing.

Features of HDFS


- **Fault tolerant** — data is duplicated across multiple DataNodes (default replication factor: 3).
- **Scalable** — reads/writes go directly to DataNodes, so capacity scales with the number of DataNodes.
- **Space** — need more storage? Add more DataNodes and rebalance.
- **Industry standard** — HBase, MapReduce, and others build on HDFS.

Write-once-read-many semantics: you can add files and delete files, but you **cannot** update/edit an existing file's contents.

Data Organization: Blocks and Replicas

- Every file written to HDFS is split into **blocks**.
- Each block is stored on one or more DataNodes; each copy is a **replica**.
- Default **block placement policy**:
 - i. First replica → the local node.
 - ii. Second replica → a different rack.
 - iii. Third replica → the same rack as the second.

Example: Splitting a File Into Blocks



Block size 128MB, file size 400MB → 4 blocks: three full 128MB blocks, plus a 16MB block. Note the **whole** 128MB is *allocated* for that last block even though only 16MB is used.

From Blocks to Mappers

HDFS block size directly sets the *default* number of mappers a MapReduce job gets — one map task per block (the same "one mapper per partition" idea from


[04_word_count_in_mapreduce.md](#)'s

"How Many Mappers Are Needed?"):

```
File size = 100 GB, Block size = 128 MB  
100 GB / 128 MB  $\approx$  800 blocks  $\rightarrow$  ~800 mappers by default
```

Smaller blocks $\rightarrow$ more, smaller mappers (more parallelism, more scheduling overhead). Larger blocks $\rightarrow$ fewer, bigger mappers (less overhead, less parallelism). This is *why* block size is configurable per job, not fixed forever at the cluster default.


Example: NameNode Metadata → DataNode Blocks



Example: NameNode Metadata → DataNode Blocks (Continued)

The NameNode's metadata is just a map from filename to block IDs
(`"/user/hadoop/input.txt": [1,2]`) — the actual block *contents*
live on the DataNodes, often with each block duplicated across more
than one node.

Replication in Action



With a replication factor of 3, each of the file's 5 blocks lands on 3 different DataNodes. **What happens if one DataNode fails?** Every block it held still has at least 2 other copies elsewhere — no data lost, and the NameNode directs re-replication to restore the count.

Higher Replication = Less Usable Storage, More Cost

5-node cluster: 1 Master (no data) + 4 DataNodes, 50TB each
(200TB raw):

Replication factor	Usable storage	Notes
1	$200 / 1 = 200\text{TB}$	no redundancy
2	$200 / 2 = 100\text{TB}$	
3	$200 / 3 \approx 67\text{TB}$	HDFS default
5	$200 / 5 = 40\text{TB}$	160TB of raw disk is <i>unusable</i> for this

More copies buys more fault tolerance, at a direct, linear cost in usable capacity.


Fault Tolerance: How Many Nodes Can Fail?

With replication factor N (and N DataNodes), $N - 1$ nodes can safely fail without losing data.

Example: replication factor 3 $\rightarrow$ 2 nodes can safely fail.

Pushing replication all the way up to N (every node has a full copy) tolerates the most failures — but at the storage cost shown on the previous slide. "Very expensive" is the exact trade-off, not a free lunch.


Read Operation in HDFS



Read Operation in HDFS (Continued)

The client asks the NameNode for block locations, then reads the actual block *data* directly from the DataNodes — the NameNode is never on the data path itself, only the metadata path.

Write Operation in HDFS



Write Operation in HDFS (Continued)

The client asks the NameNode to `create()` the file, then streams data packets through a **pipeline of DataNodes** (each acknowledging the previous one) before telling the NameNode the write is `complete()`.

HDFS Security

- **Authentication:** `simple` (trusts the OS username — insecure) or `kerberos` (ticket-based), set via `hadoop.security.authentication`.
- **File/directory permissions** follow POSIX: read/write/execute, owner/group/mode — enabled by default.
- **ACLs** handle permissions that don't fit the natural user/group hierarchy, enabled via `dfs.namenode.acls.enabled`.

Configuration Essentials

Defaults: block size **256MB**, replication factor **3**, Web UI port **50070**. Set in `hdfs-site.xml` :

```
<property>
  <name>dfs.blocksize</name>
  <value>268435456</value>
</property>
<property>
  <name>dfs.replication</name>
  <value>3</value>
</property>
```

Interfaces to HDFS

Java API (`DistributedFileSystem`), a C wrapper (`libhdfs`), HTTP, WebDAV — and the **shell**, which is usually the simplest and most familiar way in.

```
hdfs dfs      – filesystem commands (user)
hdfs fsck     – filesystem checking (user)
hdfs dfsadmin – administration commands
```


Shell Command Cheat Sheet: Files

What	Command
List a directory	<code>hdfs dfs -ls /path</code>
Disk usage	<code>hdfs dfs -du -h /path</code>
Copy local → HDFS	<code>hdfs dfs -copyFromLocal local.csv /data/</code>
Copy HDFS → local	<code>hdfs dfs -copyToLocal /data/f.csv f.csv</code>
Remove a file	<code>hdfs dfs -rm /path/file</code>
List ACLs	<code>hdfs dfs -getfacl /path/file</code>

Shell Command Cheat Sheet: Diagnostics & Admin

What	Command
Replication factor of a file	<code>hdfs dfs -stat "%r" /path/file</code>
Block locations for a file	<code>hdfs fsck /path/file -files -blocks -locations</code>
Find corrupt blocks	<code>hdfs fsck / -list-corruptfileblocks</code>
Cluster status report	<code>hdfs dfsadmin -report</code>
Rack topology	<code>hdfs dfsadmin -printTopology</code>

Summary: HDFS

HDFS is the fault-tolerant storage layer underneath Hadoop and its ecosystem (Spark, Tez, and others can all use it). Replication is what makes it fault-tolerant — data survives hardware failure — and serving reads/writes in parallel across DataNodes is what makes it fast. It's designed around **write-once, read-many-times** access.